



Assignment 3

Be sure to mark each problem # properly and your student ID (last 4 digits) shows up - no names, remember to number your pages. The submitted file should be PDF (preferably typed). Name your file with the last four digits of your student id followed by '-A3'. For example, if the last four digits of your ID are 1234, then the file name should be the following: 1234-A3.pdf

3_1 [30 points]. This assignment extends from Assignment-2: Q-1, which is reproduced below. Your assignment is to extend Part-b and use SVM, and provide a comparison with a discussion (note - need to do it only for 2.1 part-b).

PROGRAM

#import the dataset and make some changes

```
library(readr)
```

```
kc_weather_srt <- read_csv("C:/Users/bvkka/Desktop/ISL-Deep Medhi/kc_weather_srt.csv")
```

```
kc_weather_srt=kc_weather_srt[,2:9]
```

```
> kc_weather_srt
# A tibble: 366 x 8
  Temp.F Dew_Point.F Humidity.percentage Sea_Level_Press.in Visibility.mi Wind.mph Precip.in Events
  <int>    <int>          <int>          <dbl>         <int>    <int>    <dbl> <chr>
1     26         12           73         30.19           5         9     0.03  Snow
2     31         18           68         29.95           7        11     0.01  Snow
3     10          1           63         30.24           5        14     0.02  Snow
4     38         35           90         29.70           6         5     0.00  Rain
5     40         30           75         29.80           9         7     0.00  Rain
6     49         29           51         29.64          10        10     0.00  Rain
7     36         19           45         30.02          10         9     0.00  Rain
8     29         11           48         30.14          10         8     0.00  Rain
9     26          2           38         30.13          10        13     0.00  Snow
10     13        -3           46         30.37          10        12     0.00  Snow
# ... with 356 more rows
```

#first make the response column to 0-snow, 1-rain and 2-rain_thunderstorm

```
#install.packages("plyr")
```

```
library(plyr)
kc_weather_srt$Events <- revalue(kc_weather_srt$Events,c("Snow"=1))
kc_weather_srt$Events <- revalue(kc_weather_srt$Events,c("Rain"=0))
kc_weather_srt$Events <- revalue(kc_weather_srt$Events,c("Rain_Thunderstorm"=2))
```

#small changes to Events column , making it to numeric from character

```
kc_weather_srt$Events<-as.numeric(as.character(kc_weather_srt$Events))
```

```
> kc_weather_srt$Events
[1] 1 1 1 0 0 0 0 0 1 1 1 1 1 1 1 0 2 0 1 1 1 1 0 2 1 1 0 2 2 2 2 0 0 2 2 1 0 0 0 2 2 0 0 0 0 0 2 0 2 0 0 0 0 2 2 0 0 0 2 2 2 2 0 0
[66] 0 2 0 0 2 2 2 2 0 2 0 2 2 2 0 0 2 0 2 2 2 0 0 2 2 2 2 0 2 2 2 0 0 2 0 0 0 0 0 2 2 0 0 0 0 0 0 0 0 1 1 1 1 0 0 0 0 0 0 0 0 1 1
[131] 0 0 0 1 1 1 1 1 0 1 0 0 0 0 1 1 1 1 1 1 1 0 0 1 1 1 1 0 0 0 0 2 0 0 0 0 2 0 2 2 2 0 0 2 0 0 0 2 0 0 0 2 2 2 2 2 0 0 0 2 2 0
[196] 0 0 2 0 2 2 2 0 2 2 2 2 0 0 0 2 2 0 2 0 2 0 2 0 0 0 0 2 2 2 2 0 2 2 2 0 2 2 2 0 2 2 2 0 0 0 0 0 0 0 0 2 2 0 0 0 2 0 0 0
[261] 0 0 0 0 2 0 0 0 0 0 1 1 1 0 2 1 1 0 1 2 0 0 2 0 1 2 0 0 0 0 0 2 0 0 0 0 2 2 2 0 2 2 2 2 0 0 0 0 0 2 0 0 2 2 0 0 2 2 2 2 0 2 2
[326] 2 2 2 2 2 2 2 0 2 2 2 2 2 2 2 0 2 2 2 2 2 0 0 2 0 2 2 0 0 2 0 2 2 0 0 0 2 0 0 2 0 0 1 1 0
> |
```

#replications

```
rep=100
```

newly added

```
accuracy1=dim(rep)
```

```
precision_snow1=dim(rep)
```

```
precision_rain1=dim(rep)
```

```
precision_rainThunderstorm1=dim(rep)
```

```
recall_snow1=dim(rep)
```

```
recall_rain1=dim(rep)
```

```
recall_rainThunderstorm1=dim(rep)
```

#splitting the dataset into training and test sets, also install caTools packages

```
#install.packages('caTools')
```

```
library(caTools)
```

```
set.seed(123)
```

```
for(k in 1:rep)
```




```
{
```

```
  split=sample.split(kc_weather_srt$Events,SplitRatio = 0.7923)
```

```
  training_set=subset(kc_weather_srt,split==TRUE)
```

```
  test_set=subset(kc_weather_srt,split==FALSE)
```

Data

▶ kc_weather_srt	366 obs. of 8 variables	
▶ test_set	76 obs. of 8 variables	
▶ training_set	290 obs. of 8 variables	

******SVM******

#fitting SVM to the training set

```
#install.packages('e1071')
```

```
library(e1071)
```

```
classifier=svm(Events~.,data=training_set,type='C-  
classification',kernel="radial",cost=1,gamma=0.04545455,coef.0=0,epsilon=0.1)
```

```
> classifier
```

```
Call:
```

```
svm(formula = Events ~ ., data = training_set, type = "C-classification", kernel = "radial", cost = 1, gamma = 0.04545455,  
     coef.0 = 0, epsilon = 0.1)
```

```
Parameters:
```

```
  SVM-Type:  C-Classification  
 SVM-Kernel:  radial  
      cost:   1  
     gamma:  0.04545455
```

```
Number of Support Vectors: 195
```

```
y_pred1=predict(classifier,newdata = test_set[-8])
```

```
#making the confusion matrix
```

```
cm1=table(test_set$Events,y_pred1)
```

```
> cm1
```

```
  y_pred1  
    0  1  2  
0 24  0 13  
1  1  9  0  
2  7  0 22
```

```
> |
```

```
#calculating the accuracy
```

```
accuracy1[k]=mean(y_pred1==test_set$Events)
```

```
#Precision of rain, rain_thunderstorm and snow results
```

```
precision1=precision1<-diag(cm1)/colSums(cm1)
```

```
precision_rainThunderstorm1[k]=precision1[3]  
precision_snow1[k]=precision1[2]  
precision_rain1[k]=precision1[1]
```

#Recall of rain, rain_thunderstorm and snow results

```
recall1=recall1<-diag(cm1/rowSums(cm1))  
recall_rainThunderstorm1[k]=recall1[3]  
recall_snow1[k]=recall1[2]  
recall_rain1[k]=recall1[1]
```

```
}
```

#Calculating the end results using mean

```
mean(accuracy1)  
mean(precision_rain1)  
mean(precision_rainThunderstorm1)  
mean(precision_snow1)  
  
mean(recall_rain1)  
mean(recall_rainThunderstorm1)  
mean(recall_snow1)
```

SVM Radial Kernel Results

```
> mean(accuracy1)
[1] 0.7736842
> mean(precision_rain1)
[1] 0.7938852
> mean(precision_rainThunderstorm1)
[1] 0.7231856
> mean(precision_snow1)
[1] 0.8981612
> mean(recall_rain1)
[1] 0.7278378
> mean(recall_rainThunderstorm1)
[1] 0.7931034
> mean(recall_snow1)
[1] 0.887
> |
```

RESULTS:

I also changed the tuning parameters under SVM tuning to see the best results. I have used Kernels like linear, radial and sigmoid with different cost and gamma parameters. We see some differences.

SVM Linear Results

```
.6973684
[64] 0.7631579 0.7631579 0.8157895 0.7631579 0.7763158 0.7763158 0
.6973684
[71] 0.7763158 0.8026316 0.7631579 0.7631579 0.7631579 0.7500000 0
.8157895
[78] 0.7631579 0.7763158 0.7631579 0.8421053 0.8684211 0.7500000 0
.7763158
[85] 0.7631579 0.7763158 0.6973684 0.7763158 0.7763158 0.6710526 0
.7894737
[92] 0.7894737 0.7763158 0.7236842 0.7368421 0.7763158 0.8157895 0
.7763158
[99] 0.7368421 0.7368421
> mean(accuracy1)
[1] 0.7638158
> mean(precision_rain1)
[1] 0.7856708
> mean(precision_rainThunderstorm1)
[1] 0.7071105
> mean(precision_snow1)
[1] 0.9012634
> mean(recall_rain1)
[1] 0.7132432
> mean(recall_rainThunderstorm1)
[1] 0.7817241
> mean(recall_snow1)
[1] 0.899
> |
```

SVM RADIAL WITH GAMMA =0 AND COST =1 RESULTS

```
+ precision_snow1[k]=precision1[2]
+ precision_rain1[k]=precision1[1]
+
+
+ recall1=recall1<-diag(cm1/rowSums(cm1))
+ recall_rainThunderstorm1[k]=recall1[3]
+ recall_snow1[k]=recall1[2]
+ recall_rain1[k]=recall1[1]
+
+
+
+ }
> mean(accuracy1)
[1] 0.7765789
> mean(precision_rain1)
[1] 0.7504617
> mean(precision_rainThunderstorm1)
[1] 0.7881519
> mean(precision_snow1)
[1] 0.8908958
> mean(recall_rain1)
[1] 0.8143243
> mean(recall_rainThunderstorm1)
[1] 0.7182759
> mean(recall_snow1)
[1] 0.806
> |
```


Model	Tuning Parameters	Accuracy	Precision Snow	Precision Rain	Precision Rain Thunderstorm	Recall Snow	Recall Rain	Recall Thunderstorm
SVM	kernel="radial",cost=1, gamma=0.04545455,coef.0=0,epsilon=0.1	0.7736842	0.8981612	0.7938852	0.7231856	0.887	0.7278378	0.7931034
SVM	kernel="linear"	0.7638158	0.9012634	0.7856708	0.7071105	0.899	0.7132432	0.7817241
SVM	kernel="radial",cost=1,gamma=0	0.7765789	0.8908958	0.7504617	0.7881519	0.806	0.8143243	0.7182759
SVM	kernel="sigmoid",cost=1, gamma=0.04545455,coef.0=0,epsilon=0.1	0.7515789	0.8825927	0.7829367	0.6924296	0.868	0.68	0.8027586

Comparing to the other models using in Assignment 2

Model	Accuracy	Precision Snow	Precision Rain	Precision Rain Thunderstorm	Recall Snow	Recall Rain	Recall ThunderStorm
LDA	0.9026316	0.6407459	0.9115871	0.9906168	0.911675	0.902705	0.9906152

QDA	0.7389474	0.945	0.697027	0.7213793	0.7950919	0.7514844	0.7115056
KNN (K=5)	0.745	0.895	0.7305405	0.7117241	0.9098042	0.7444542	0.701065

Discussion Note:

1. From Accuracy Results, we see that SVM model performs better than QDA and KNN, but LDA outperforms SVM too.
2. From Precision of Snow Results, SVM does better than LDA and KNN
3. From Precision of Rain Results, SVM does better than QDA and KNN, but less than LDA
4. From Precision of thunderstorm Results, SVM does better than QDA and KNN, but less than LDA.
5. From Recall of Snow Results, SVM does better than QDA and KNN
6. From Recall of Rain Results, SVM does better than QDA and KNN, but less than LDA
7. From Recall of thunderstorm Results, SVM does better than QDA and KNN, but less than LDA.
8. So, overall if we compare performance with respect to classifiers, LDA>SVM>KNN>QDA.

3_2. [45 points] Consider the time series on Milk production data [milk-production\(1\).csv](#)

it shows cow milk production per pound from 1962 to 1975.

- Try at least three different values for window size with simple moving average (SMA) for forecasting
- Apply exponential moving average using HoltWinters for forecasting
- For the above, discuss how the forecasting differs in terms of MAD and MFE and why one approach or the other is better.

PROGRAM

#import the dataset and make some changes

```
library(readr)
```

```
milk_production <- read_csv("C:/Users/bvkka/Desktop/ISL-Deep Medhi/assignment3/milk-production(1).csv")
```

```
> milk_production
```

```
# A tibble: 168 x 3
```

```
  Month `Pounds_per_Cow` \r\n"1962-01` `589`  
  <chr>                <int> <chr>  
1 1962-01                589 <NA>  
2 1962-02                561 <NA>  
3 1962-03                640 <NA>  
4 1962-04                656 <NA>  
5 1962-05                727 <NA>  
6 1962-06                697 <NA>  
7 1962-07                640 <NA>  
8 1962-08                599 <NA>  
9 1962-09                568 <NA>  
10 1962-10               577 <NA>
```

```
# ... with 158 more rows
```

```
> |
```

```
milk_production<-milk_production[,2]
```

```
> milk_production
# A tibble: 168 x 1
  `Pounds_per_Cow` \r\n"1962-01`
      <int>
1         589
2         561
3         640
4         656
5         727
6         697
7         640
8         599
9         568
10        577
# ... with 158 more rows
> |
```

```
milk_production_timeseries<-ts(milk_production)
```

#contains monthly milk productions for January 1970-Decemeber 1983


```
mp_TS<-ts(milk_production,frequency = 12,start=c(1970,1))
```

```
> mp_TS
      Jan  Feb  Mar  Apr  May  Jun  Jul  Aug  Sep  Oct  Nov  Dec
1970  589  561  640  656  727  697  640  599  568  577  553  582
1971  600  566  653  673  742  716  660  617  583  587  565  598
1972  628  618  688  705  770  736  678  639  604  611  594  634
1973  658  622  709  722  782  756  702  653  615  621  602  635
1974  677  635  736  755  811  798  735  697  661  667  645  688
1975  713  667  762  784  837  817  767  722  681  687  660  698
1976  717  696  775  796  858  826  783  740  701  706  677  711
1977  734  690  785  805  871  845  801  764  725  723  690  734
1978  750  707  807  824  886  859  819  783  740  747  711  751
1979  804  756  860  878  942  913  869  834  790  800  763  800
1980  826  799  890  900  961  935  894  855  809  810  766  805
1981  821  773  883  898  957  924  881  837  784  791  760  802
1982  828  778  889  902  969  947  908  867  815  812  773  813
1983  834  782  892  903  966  937  896  858  817  827  797  843
> |
```

```
#plotting time series
```

```
plot.ts(milk_production_TS)  
lines(mp_TS,col="blue")
```

Plot of Timeseries



****Question #2 – Part A****

****Simple Moving Average(SMA)***# ->it is used to smooth time series data*

```
#install.packages("TTR")
```


```
library("TTR")
```

```
mp_TS5<-SMA(milk_production_TS,n=5)
```

```
plot.ts(mp_TS5)
```


```
lines(mp_TS5,col="purple")
```

Simple Moving Average: Window size 5



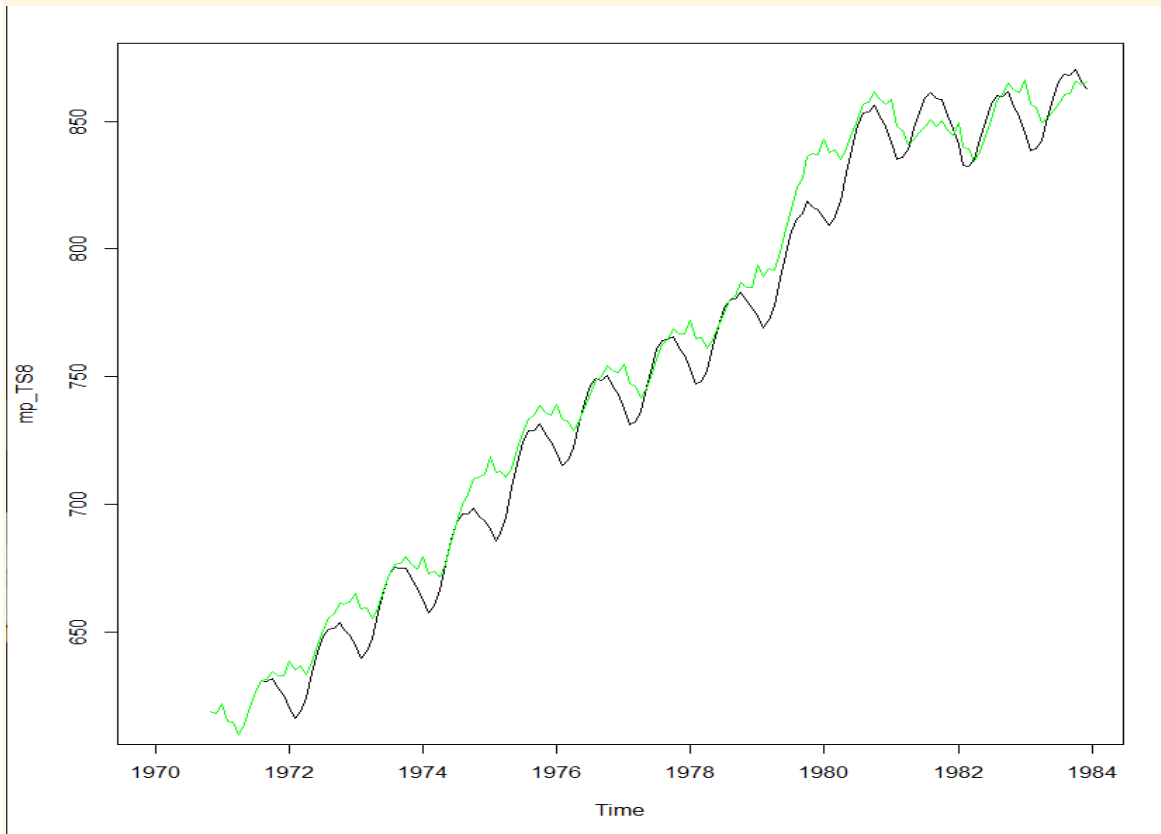
```
mp_TS5<-SMA(milk_production_TS,n=11)  
plot.ts(mp_TS11)  
lines(mp_TS11,col="red")
```

Simple Moving Average: Window size 11



```
mp_TS20<-SMA(milk_production_TS,n=20)  
plot.ts(mp_TS20)  
lines(mp_TS20,col="green")
```

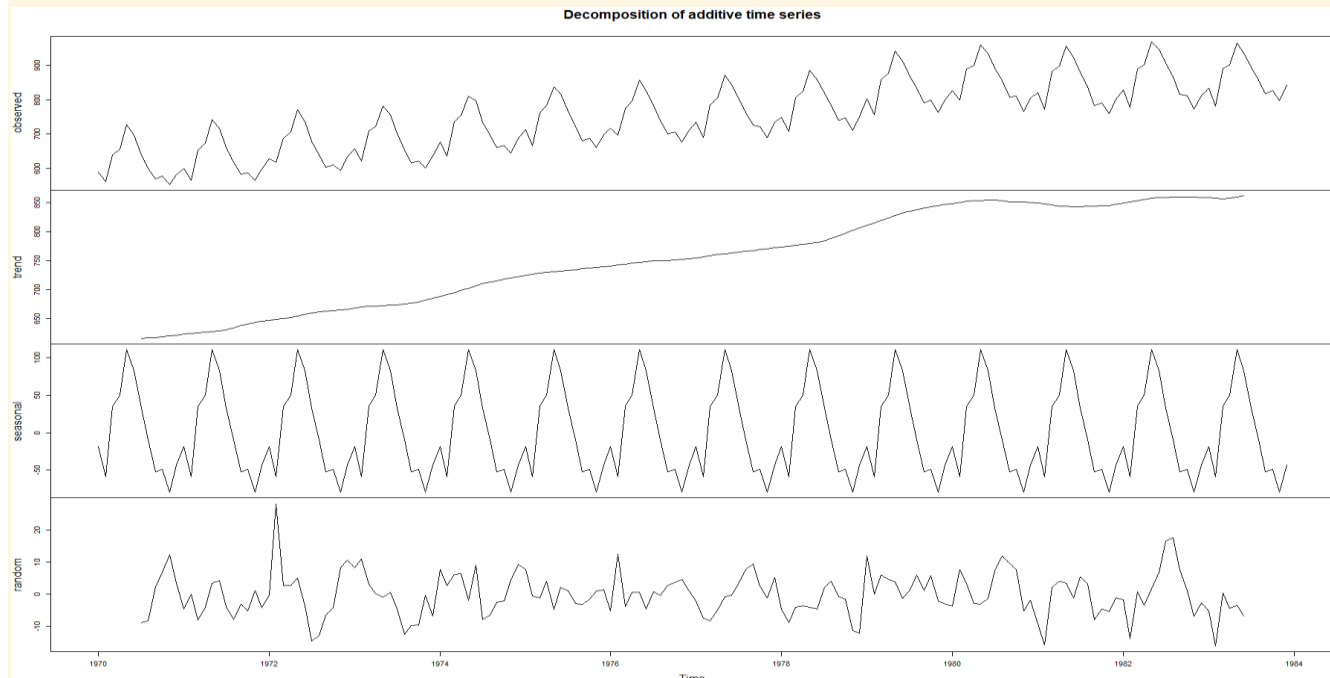
Simple Moving Average: Window size 20



#To estimate the trend component and seasonal component of a seasonal time series that can be described using an additive model, we can use the “decompose ()” function in R. This function estimates the trend, seasonal, and irregular components of a time series that can be described using an additive model###

```
mp_decompose=decompose(milk_production_TS)
```

plot(mp_decompose) #The plot above shows the original time series (top), the estimated trend component (second from top), the estimated seasonal component (third from top), and the estimated irregular component (bottom)##



Discussion and observations:


- The moving average model uses the last t periods to predict demand in period $t+1$.
- SMA is an arithmetic moving average calculated by adding the actual forecasts for many time periods and then dividing this total by the number of time periods.
- In the above program, simple moving average of milk-production dataset has been calculated, with three different windows, 5,11,20. I chose this window sizes because,
- When window size is 5, the graph is changing, but we can see that it is changing in in equally distributed patterns. Whereas when window size is 11, the graph is increasing but with some non-linearity. When the window size is 20, the graph is almost linearly increasing.

Question #2 – Part B

*****Forecasts suing Exponential Smoothing****###*

```
mp_exp=HoltWinters(mp_TS,beta=FALSE,gamma = FALSE)
plot(mp_exp) #The plot shows the original time series in black, and the forecasts as a red line. The time series of forecasts is much smoother than the time series of the original data here.
lines(mp_exp$fitted[,1],col="red")
```

Holt-Winters filtering



mp_exp

```
> mp_exp
Holt-Winters exponential smoothing with trend and additive seasonal component.

Call:
HoltWinters(x = mp_TS)

Smoothing parameters:
  alpha: 0.68933
  beta : 0
  gamma: 0.8362592

Coefficients:
      [,1]
a  885.775547
b    1.278118
s1 -16.743296
s2 -59.730034
s3  47.492731
s4  56.203890
s5 115.537545
s6  84.554817
s7  39.580306
s8  -4.702033
s9 -54.554684
s10 -51.582594
s11 -85.953466
s12 -42.907363
```

```
#install.packages("TTR")
library(forecast)
forecast_holtwinter=forecast(mp_exp)
forecast_holtwinter
```


```
> forecast_holtwinter
      Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
Jan 1984      842.997 784.7725 901.2214 753.9504 932.0436
Feb 1984      842.997 760.6579 925.3360 717.0702 968.9237
Mar 1984      842.997 742.1537 943.8402 688.7705 997.2234
Apr 1984      842.997 726.5539 959.4401 664.9126 1021.0813
May 1984      842.997 712.8100 973.1839 643.8932 1042.1007
Jun 1984      842.997 700.3846 985.6093 624.8902 1061.1037
Jul 1984      842.997 688.9583 997.0356 607.4152 1078.5788
Aug 1984      842.997 678.3229 1007.6710 591.1497 1094.8442
Sep 1984      842.997 668.3339 1017.6600 575.8729 1110.1211
Oct 1984      842.997 658.8861 1027.1079 561.4236 1124.5703
Nov 1984      842.997 649.8999 1036.0940 547.6806 1138.3134
Dec 1984      842.997 641.3138 1044.6801 534.5492 1151.4447
Jan 1985      842.997 633.0786 1052.9154 521.9545 1164.0394
Feb 1985      842.997 625.1544 1060.8395 509.8356 1176.1584
Mar 1985      842.997 617.5086 1068.4853 498.1423 1187.8517
Apr 1985      842.997 610.1136 1075.8803 486.8326 1199.1613
May 1985      842.997 602.9464 1083.0475 475.8713 1210.1226
Jun 1985      842.997 595.9870 1090.0069 465.2278 1220.7661
Jul 1985      842.997 589.2184 1096.7755 454.8761 1231.1178
Aug 1985      842.997 582.6257 1103.3682 444.7935 1241.2004
Sep 1985      842.997 576.1959 1109.7981 434.9599 1251.0340
Oct 1985      842.997 569.9174 1116.0766 425.3578 1260.6362
Nov 1985      842.997 563.7800 1122.2139 415.9715 1270.0224
Dec 1985      842.997 557.7747 1128.2192 406.7871 1279.2068
```

```
Accuracy(forecast_holt_winters)
```

```
> accuracy(forecast_holtwinter)
              ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set 1.52104 45.32207 39.0064 0.03434009 5.159602 1.832279 0.02263573
> |
```

```
plot(forecast_holt_winters)
```

Forecasts from HoltWinters



Discussion and observations:

- The main idea is an exponential smoothing is that the prediction mostly depends on most recent observation and on the error of the latest forecast.
- If the time series can be described using an additive model with increasing or decreasing trend and seasonality, Holt-Winters exponential smoothing to make short-term forecasts.
- Smoothing is controlled by three parameters: alpha, beta, and gamma, for the estimates of the level, slope b of the trend component, and the seasonal component, respectively, at the current time point. The coefficients alpha, beta and gamma, usually ranges between 0 and 1.

Holt-Winters exponential smoothing with trend and additive seasonal component.

Call:

```
HoltWinters(x = mp_TS)
```

Smoothing parameters:

alpha: 0.68933

beta : 0

gamma: 0.8362592

-
- In the above forecast, α is 0.68933 indicating that the estimate of the level at the current time is based upon observations in the more distant past as well as some recent observations.
- The value of beta is 0, which indicates that the estimate of the slope b of the trend component is not updated over the time series, and instead is set equal to its initial value. This makes a conclusion that as the level changes over the time series, slope b of the trend component remains almost constant.
- In contrast, the value of gamma (0.8362592) is high, indicating that the estimate of the seasonal component at the current time point is just based upon very recent observations.

***Question #2 – Part C**

Discuss how the forecasting differs in terms of MAD and MFE and##
##why one approach or the other is better##

```
count=168
x=mean(mp_TS3[-(1:7)])
for (k in 8:count) {
  mean_abs_dev3=mean(abs(mp_TS3[k]-x))
  mfe3=mean(mp_TS3[k]-x)
  mad3=mad(mp_TS3[k], centre, constant = 1.4826, na.rm = FALSE, low = FALSE, high = FALSE)
}
y=mean(mp_TS5[-(1:12)])
for (k in 13:count) {
  mean_abs_dev4=mean(abs(mp_TS5[k]-y))
  mfe4=mean(mp_TS5[k]-y)
  mad4=mad(mp_TS5[k], centre, constant = 1.4826, na.rm = FALSE, low = FALSE, high = FALSE)
}
z=mean(mp_TS8[-(1:23)])
for (k in 24:count) {
  mean_abs_dev5=mean(abs(mp_TS8[k]-z))
  mfe5=mean(mp_TS8[k]-z)
  mad(mp_TS8[k], centre, constant = 1.4826, na.rm = FALSE, low = FALSE, high = FALSE)
}
```

```
> mfe3
[1] -93.8795
> mfe4
[1] -135.6247
> mfe5
[1] -135.1197
> mean_abs_dev3
[1] 93.8795
> mean_abs_dev4
[1] 135.6247
> mean_abs_dev5
[1] 135.1197
> |
```



```
> accuracy(forecast_holtwinter)
              ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
Training set 1.52104 45.32207 39.0064 0.03434009 5.159602 1.832279 0.02263573
> |
```

Observations:

- According to the above values of MAD and MFE of Simple moving average (three different windows) and MAE of Holt-winter forecasting, I find Holt-winters exponential smoothing and forecasting more precise than SMA.